

Intro IoT Machine Learning Project 1

Jaskin Kabir

Student Id: 801186717

GitHub:

https://github.com/holleman-courses/proj1-awesome_team/

March 2025

1 Introduction

The purpose of this project was to develop a machine learning model capable of recognizing pencils and deploy this model to the Arduino Nano 33 BLE Sense. This project aims to demonstrate the feasibility of deploying a machine learning model on embedded hardware with limited computational resources. By focusing on a simple object like a pencil, the project provides a manageable starting point for exploring the challenges of embedded machine learning, such as memory constraints, inference speed, and accuracy.

The intention is to produce a system that is not only sufficiently accurate but also fast enough to be used in real-time applications. The project involves designing a neural network architecture that balances performance and resource efficiency, constructing a dataset tailored to the target object, and training the model to achieve optimal accuracy. The upcoming sections will delve into the details of the model architecture, the dataset preparation process, the training methodology, and the results obtained. These sections will also discuss the challenges encountered, and the lessons learned during the development and deployment of the system.

2 Model Architecture

For image classification on embedded hardware, the most widely used paradigm is the convolutional neural network, which was used in this project. As for the structure of the network itself, a choice was made between a wide and short network or a narrow and deep network. A wide and short network would have few convolutional layers with each layer consisting of many convolutional filters, whereas the deep and narrow network would have many layers with fewer filters in each layer. To allow for the greatest abstraction from the input images, the narrow and deep network was chosen. To resolve the issue of vanishing gradients, skip connections were added between every two convolutional layers after the initial layer.

Each residual block consists of two convolutional layers with batch normalization and relu activation functions between them. The structure of the resblock can be seen in Figure 1. The overall network topology is shown in Figure 2. The initial convolution uses a stride of 4 to reduce the size of the input image so the microprocessor could process it. After the initial convolution, the stride is set to 1 for the rest of the network. The size of the network was mainly constrained by the runtime, as the required tensor arena size had to fit within the small RAM of the Arduino. The parameter count was kept low enough to allow for faster training and inference times, but large enough to accurately classify images. To keep the required arena size low, the maximum number of filters in each layer was kept at 128, and on either side of the 128 filter layers is a 64 filter layer.

3 Data and Training

The dataset used to train and evaluate the model was a set of 896 grayscale images taken using the Arduino Nano and the OV767X camera included in the kit. The images were taken at a 176x144 resolution and manually labeled. The images include pictures on various surfaces like chairs, jackets, bedsheets and tables. 459 images contained a pencil and 437 did not. Of these negative images, some contained pictures of screwdrivers or wrenches so that the model would not learn to classify straight objects as pencils.

The images weren't resized or preprocessed in any way, other than random rotation and flipping for data augmentation. The images were not resized before training so that rescaling logic wouldn't need to be implemented on the Arduino. The dataset was split into a training set of 70% and a validation set of 30%. The model was trained for 100 epochs, with a batch size of 32. The model was trained using the Adam optimizer with a learning rate of 0.005. As seen in Figure 3, the model suffered from significant overfitting. Additionally, the validation curves are very unstable. This is likely due to the small validation dataset.

4 Results

In software, the validation accuracy was 88%. When deployed to the Arduino, the system did respond to the target object, but false positives were easy to induce. When pointed at a granite countertop with a striped pattern, the stripes were likely mistaken for pencils. Additionally, the model tends to classify wires as pencils as well. This may have been alleviated if more pencil-like objects were included in the negative dataset.

Inference takes 4.883 seconds per image. The frame rate of the camera was set to 5fps, as the dataset images were captured at that frame rate to reduce blurriness.

5 Summary Table

Metric	Value
Validation Accuracy	87.73%
Training Accuracy	95.07%
Inference Time	4.883s
False Rejection Rate	16.33%
False Positive Rate	7.38%
Parameter Count	112634
Input Shape	176x144x1
MACs	40,278,240
Frame Rate	5fps

Table 1: Summary of results

6 Discussion

Overall, the system developed is a good proof of concept. It successfully deploys a resnet based object detection model that does respond to the target class. However, the inference time is very slow, and the accuracy leaves much to be desired. As the system takes almost 5 seconds to process images, it would not be useful in real-time applications. The biggest challenge was handling the small dataset. Augmentation can only go so far, and the lack of validation or test data led to unstable results. Thus, the first step that should be taken to improve the system would be to collect more data. To improve inference time, a much shallower model should be used with fewer parameters.

7 Conclusion

Through this project, I learned how to deploy an object detection model on an embedded system. I learned how to construct a dataset and label it. I also learned how to design a model's architecture to fit within the constraints of the hardware. I learned how to train a model and evaluate its performance. I also learned how to deploy a model to an embedded system and measure its performance. This was the biggest challenge, as many times I would build a model, but it would fail upon deployment due to memory constraints.




Figure 1: Residual block




Figure 2: Network topology




Figure 3: Loss and accuracy curves